

My grades for Homework 3

Q1

25 / 25

Refer to the assignment PDF.

CS 466/666: Algorithm Design and Analysis University of Waterloo Fall 2023

Homework 3 Due: Monday, November 13, 2023

Name: William Gao, Ades Li, Anthony Tieu

Problem 1. Suppose we have a bipartite graph $G = (L, R, E)$ and our goal is to find a minimum size set of edges F such that every vertex in $L \cup R$ is incident on at least one of the edges in F . Design a polynomial time algorithm for this problem. (25 points)

Hint: You can use whatever algorithm you like, but perhaps a simple way is to write a standard LP relaxation for this problem and then show that for bipartite graphs, the integrality gap of this LP is 1 (using an approach you have already seen in Lecture 8).

Solution. Solution.

ILP:

- We define an integral variable $y_e \in \{0,1\}$ for each edge $e \in E$ (where $y_e = 1$ if $e \in F$ and $y_e = 0$ otherwise).
- Then we require that
$$\forall s \in L : \sum_{e \in E} y_e \geq 1$$

$$\forall r \in R : \sum_{e \in E} y_e \geq 1$$
- We define our minimization problem as
$$\min \sum_{e \in E} y_e$$
- Thus, **ILP** is defined as:
$$\begin{aligned} & \min \sum_{e \in E} y_e \\ & \text{subject to } \forall s \in L : \sum_{e \in E} y_e \geq 1 \\ & \quad \forall r \in R : \sum_{e \in E} y_e \geq 1 \\ & \quad y_e \in \{0,1\} \end{aligned}$$
✓

LP Relaxation:

The LP Relaxation is then given by:

$$\begin{aligned} & \min \sum_{e \in E} x_e \\ & \text{subject to } \forall s \in L : \sum_{e \in E} x_e \geq 1 \\ & \quad \forall r \in R : \sum_{e \in E} x_e \geq 1 \\ & \quad x_e \in [0,1] \end{aligned}$$

1 ✓

The objective value of the LP Relaxation serves as a lower bound to the ILP (given that any solution to the ILP is also a solution to the LP Relaxation).

Proposition 1:
Any feasible and optimal solution $x \in \mathbb{R}^E$ to the LP Relaxation for bipartite edge cover can be rounded to a feasible integral solution $\bar{x} \in \{0, 1\}^E$ without increasing the objective value.

Proof of Proposition 1. Fix the feasible and optimal solution $x \in \mathbb{R}^E$. If x is already integral, we are done by taking $\bar{x} = x$. Otherwise, consider the following argument. We define the support of x as $\text{supp}(x) = \{e \in E : x_e > 0\}$.

Claim 1: Let a be any leaf-node in $\text{supp}(x)$ (if a leaf-node exists) and let i be the edge incident on a . By feasibility, $\sum_{v \in N(a)} x_v \geq 1$, as each, by optimality, $x_i = 1$.

Step 1: Cycle Cancelling.

- Note: If x is the optimal solution to the LP Relaxation, then $\text{supp}(x)$ cannot contain an integral cycle (i.e. a cycle such that all edges are assigned a value $x_e = 1$). If such a cycle did exist, then by removing one edge from it, all vertices within the cycle would still be incident on at least one edge and thus the sum of the fractional values of the edges in the cycle would be smaller than that of x , a contradiction. Thus, any cycle in $\text{supp}(x)$ must be fractional (at least, in part).
- We perform Cycle Cancelling as defined in LeS8 section 2.3:
 - Find a cycle in the support of x that contains a fractional value. If no cycle can be found, then proceed to Step 2.
 - Let $C = e_1 \rightarrow e_2 \rightarrow \dots \rightarrow e_k \rightarrow e_1$ be our cycle where $e_1 \in C$ is the edge with the smallest fractional value in the cycle. Let $\delta = x_{e_1}$.
 - Update x as follows: $x_{e_1} \leftarrow x_{e_1} - \delta, x_{e_2} \leftarrow x_{e_2} + \delta, x_{e_3} \leftarrow x_{e_3} - \delta, \dots$
- This guarantees:
 - At least one edge is removed from $\text{supp}(x)$ (since $x_{e_1} > 0$).
 - Since G is bipartite, its cycles must be even, as such $\sum_{e \in C} x_e$ remains the same for all $b \in L \cup R$ (we are increasing the value of one edge of b by δ and decreasing the value of another edge of b by the same value). Feasibility is retained.
 - **Claim 2:** Let a be any leaf-node produced by Cycle Cancelling and let i be the edge incident on a .
Proof of Claim 2. Since feasibility is retained, then $\sum_{v \in N(a)} x_v \geq 1$, thus $x_i \geq 1$. Since optimality is retained, then $x_i = 1$ (otherwise, if $x_i > 1$, then we could reduce $x_i = 1$ and obtain a solution with a higher objective value). As such, Cycle Cancelling is $O(E^2)$ (polynomial time).

Step 2: Rounding Forests.

- No cycles remain in the support of x , so $\text{supp}(x)$ is a forest. There must always be a leaf-node (unless the forest is empty).

2

• Let u be a leaf-node in $\text{supp}(x)$, let v be its parent, and let $i \in \{u, v\}$ be the edge between them.

• By Claim 2, x_i is integer (as v is a leaf-node).

• Suppose that v has child-nodes (other than u). Let w be an arbitrary child-node of v (where $w \neq u$) and let $j \in \{v, w\}$ be the edge between them. Since there are no cycles in the support of x , then each child w of v can be handled independently.

- If w is not a leaf-node, then we distribute the x -value of j amongst the other edges incident on w and we set $x_j = 0$. We note that this does not affect $\sum_{v \in N(u)} x_v$ and it can only serve to possibly increase the sum of the x -values of the neighbours v of u (with the exception of v).
- This guarantees:
 - * $\sum_{v \in N(u)} x_v$ increases after this step for all $v \in L \cup R$ (with the exception of v , but we note that v is a leaf-node, so $\sum_{v \in N(u)} x_v \geq 1$ given that $x_u = 1$). Feasibility is retained.
 - * $\sum_{v \in N(u)} x_v$ remains the same. Optimality is retained.
- **Claim 3:** Let a be any leaf-node produced by this process and let i be the edge incident on a .
Proof of Claim 3. Since feasibility is retained, then $\sum_{v \in N(a)} x_v \geq 1$, thus $x_i \geq 1$. Since optimality is retained, then there must exist some method of distributing the x -value of j amongst the other edges incident on w such that $x_i = 1$ (otherwise, if such a method did not exist and $x_i < 1$, then we could reduce $x_i = 1$ and obtain a solution that is smaller than the optimal solution, a contradiction).

• Repeat for all possible child-nodes w of v .

• Ignore u, v , and w in subsequent iterations of Rounding Forests.

• Repeat until no leaf-nodes remain.

• Note that each iteration of Rounding Forests is $O(V)$ (v has less than $|V|$ child-nodes) and since we ignore two nodes per iteration, the total number of iterations of Rounding Forests is also $O(V)$. As such, Rounding Forests is $O(V^2)$ (polynomial time).

At this point, all x -values are integer. The solution remains feasible and retains the optimality of the original LP Relaxation solution. Since the LP Relaxation provides a lower bound to the ILP, then x is an optimal solution to the ILP, as desired. ■

3

Q2

25 / 25

Refer to the assignment PDF.

Problem 2: Let $G = (V, E)$ be any graph with integer weights $w(e)$ on edges $e \in E$. For this question, we allow the edges to have negative weight. Consider the following approximation algorithm for the maximum weight matching problem, namely, finding a matching $M \subseteq E$ that maximizes $\sum_{e \in M} w(e)$:

1. If all edges in G have a non-positive weight, return $M = \emptyset$ and terminate.
2. Pick an arbitrary edge $e \in E$ with $w(e) > 0$. Create the new graph $G' := G - e$ with weights $w'(f) = w(f) - w(e)$ for every edge $f \in E$ such that e and f incident on v (basically, we are subtracting $w(e)$ from the weights of all edges incident on v in G').
3. Run the algorithm recursively on G' to obtain a matching M' . If both endpoints of the edge e are unmatched, return $M := M' \cup \{e\}$; otherwise, return $M := M'$ as the final answer.

Prove that this algorithm outputs a $(1/2)$ -approximation to the maximum weight matching problem.

(25 points)

Solution. Let's first consider an **Invariant 1** in the provided algorithm, for any G and G' where $G' := G - e$ with the updates on all the edges that are incident on e . Suppose the optimal maximum weight matching for G is Opt_G , and the optimal maximum weight matching for G' is $Opt_{G'}$. Let M and M' denote the maximum weight matching for G and G' returned by the algorithm respectively.

- If $e \notin Opt_G$:
Since there can be at most one edge in Opt_G that is incident on e , $w'(Opt_{G'}) \geq w(Opt_G) - 2w(e)$.
- If $e \in Opt_G$:
Since there can be at most 2 edges from Opt_G that are incident on e , $w'(Opt_{G'}) \geq w(Opt_G) - 2w(e)$.

So, we have **Invariant 1**: For any G' retrieved from G by step 2, $w'(Opt_{G'}) \geq w(Opt_G) - 2w(e)$.

Now, let's prove that this algorithm computes a $(1/2)$ -approximation to the maximum weight matching problem after backtracking by an induction.

Base case: If all edges in G have a non-positive weight, M will be empty. So, $w(M) = 0$. It's clear that $w(Opt_G) = 0$, so $w(M) \geq \frac{1}{2}w(Opt_G)$. \square

Then, we can define our **inductive hypothesis**. Assume that for any intermediate G_k , $w_k(M_k) \geq \frac{1}{2}w_k(Opt_{G_k})$. The remained is to prove that for the original graph G , $w(M) \geq \frac{1}{2}w(Opt_G)$.

By step 3, we know that $M := M' \cup e$ or $M := M'$ depending on whether both endpoints of the edge e are unmatched in M' :

- (1) If $M := M' \cup e$, both endpoints of e are not in M' . So, we don't need to adjust the weight of the f in M of G since none of them can be incident on e . And the only adjustment required is to add $w(e) - w(M')$:

4

$$\begin{aligned} w(M) &= w'(M') + w(e) \\ &\geq \frac{1}{2}w'(Opt_{G'}) + w(e) \quad (\text{By III}) \\ &\geq \frac{1}{2}(w(Opt_{G'}) - 2w(e)) + w(e) \quad (\text{By Invariant 1}) \\ &\geq \frac{1}{2}w(Opt_{G'}) - w(e) + w(e) \\ &= \frac{1}{2}w(Opt_{G'}) \end{aligned}$$

(2) If $M := M'$, then at least one of the endpoints of e is matched in M' . So, when we backtrack from G' to G , we update the weight of the edges $f \in M'$ that are incident on e by adding back the reduced portion of the weight $w(e)$ in Step 1. It's clear that the # of f whose weight needs to be incremented by $w(e)$ is at least 1 and at most 2. Therefore,

$w'(M') + 2w(e) \geq w(M) \geq w'(M') + w(e)$

So,

$$\begin{aligned} w(M) &\geq w'(M') + w(e) \\ &\geq \frac{1}{2}w'(Opt_{G'}) + w(e) \quad (\text{By III}) \\ &\geq \frac{1}{2}(w(Opt_{G'}) - 2w(e)) + w(e) \quad (\text{By Invariant 1}) \\ &\geq \frac{1}{2}w(Opt_{G'}) - w(e) + w(e) \\ &= \frac{1}{2}w(Opt_{G'}) \end{aligned}$$

Hence, by induction, we have the induction conclusion that this algorithm outputs a $(1/2)$ -approximation to the maximum weight matching problem. \blacksquare

5

Solution. We first prove the following claim:

Claim 1: Consider a recursive step running on graph G' with weights u' as defined in the algorithm. The matching M' returned by the algorithm in this step is a $(1/2)$ -approximation to the maximum weight matching problem on G' with weights u' .

It follows from the claim that the algorithm outputs a $(1/2)$ -approximation by applying the claim on the original graph G with weights u .

Proof. Base case: All edges have non-negative weight. The maximum matching has weight 0 and the algorithm returns \emptyset which also has weight 0.

Inductive step: Suppose our algorithm is running on G' with weights u' . Let e be an edge with positive weight ($u'(e) > 0$). Let $\tilde{G} = G' - e$ be the new graph generated by step 2 of the algorithm with weights $\tilde{w}(f) = u'(f)$ for every edge f not incident on e and $\tilde{w}(f) = u'(f) - u'(e)$ for every edge f incident on e .

By the inductive hypothesis, the algorithm will return a matching \tilde{M} such that \tilde{M} is a $(1/2)$ -approximation of a maximum matching on \tilde{G} .

Let M' be the matching constructed in step 3 of the algorithm where $M' = \tilde{M} \cup \{e\}$ if both endpoints u, v of e are unmatched, otherwise $M' = \tilde{M}$. We wish to prove that M' is a $(1/2)$ -approximation on G' , u' .

Since $\tilde{M} \subseteq M'$, M' is a $(1/2)$ -approximation for \tilde{G} , \tilde{w} . Now, examine the weight function

$$\bar{w} = \begin{cases} u'(e) & \text{if } f \text{ is incident on } e \\ 0 & \text{otherwise} \end{cases}$$

Then, since the algorithm in step 2, subtracts $u'(e)$ from $u'(f)$ if f is incident on e to get \tilde{w} ,

$$u' = \tilde{w} + \bar{w}.$$

Now consider a maximum matching on \tilde{G} . Since there are only two vertices u, v that are incident to e and any edge that is not incident to e has weight 0 under \tilde{w} , a max matching on \tilde{G} has weight at most $2u'(e)$. Furthermore, M' must have at least one edge incident to u or v , otherwise we would have added e to M' . So M' has weight at least $\frac{1}{2}u'(e)$ under u and thus M' is a $(1/2)$ -approximation under u :

Now, to conclude the proof, we let M be a max matching under u' on G' . We show that $u'(M') \geq \frac{1}{2}u'(M)$:

$$\begin{aligned} u'(M') &= \tilde{w}(M') + \bar{w}(M') && \text{(From (1))} \\ &\geq \frac{1}{2}\tilde{w}(M) + \frac{1}{2}\bar{w}(M) && \text{(Since } M' \text{ is a } (1/2)\text{-approximation for both } \tilde{w} \text{ and } \bar{w}) \\ &= \frac{1}{2}u'(M) \end{aligned}$$

as desired. \square

6

Q3

21.5 / 25

Refer to the assignment PDF.

Problem 3. Consider the following online problem. Unfortunately, you have lost a bet to your (most untrustworthy) friend and now have to accommodate him every day, you either have to buy lunch for your group of friends or pay L dollars on that day, you can give up and end the bet at any time, or you can pay your friend to end your misery (at least in this bet). However, your friend can also decide at any day that they are done with this bet and no longer want to continue this and you have no idea on if or when they do this. You may assume in this question that L divides G to simplify the math.

1. Design a **deterministic** strategy such that the total money you spend is at most $(2 - \frac{G}{L})$ times the minimum amount you had to spend if you knew which day your friend decides to stop this bet. (12.5 points)

Solution. Solution.

Deterministic Strategy:

- (a) Pay L dollars for $\frac{G-L}{L}$ days. ✓

- (b) Then give up and pay G . ✓

Proof of Deterministic Strategy. Let d denote the number of days our unscrupulous friend persists with the bet.

Claim 1. The minimum amount we could possibly spend is either $L \cdot d$ or G . By the rules of the bet, we must either pay L dollars for d days; pay L dollars for $a \neq (0, d]$ days, then pay G dollars; or pay G dollars on the first day. Note that paying L dollars for $a \neq (0, d]$ days, then paying G dollars is never optimal since then we simply pay G dollars on the first day.

We wish to show that our strategy works for all values of d :

- Case 1: $d \leq \frac{G-L}{L}$.

Total amount we pay: $L \cdot d$ dollars.

We wish to show that paying $L \cdot d$ dollars total is optimal when $d \leq \frac{G-L}{L}$ (i.e., by **Claim 1**, we want to show that $L \cdot d \leq G$):

$$\begin{aligned} L \cdot d &\leq L \cdot \frac{G-L}{L} \\ &= G - L \\ &< G \end{aligned}$$

Since it is already optimal, it is certainly within $(2 - \frac{G}{L})$ times the optimal value.

- Case 2: $d \geq \frac{G-L}{L}$ (or $d \geq \frac{G-L}{L} + 1$)

Total amount we pay: $\frac{G-L}{L} \cdot L + G = G - L + G = 2G - L$ dollars.

Observe that G becomes the optimal value when $L \cdot d \geq G$ (or when $d \geq \frac{G}{L}$). When $d \geq \frac{G-L}{L} + 1$, we have:

$$\begin{aligned} d &\geq \frac{G-L}{L} + 1 \\ d &\geq \frac{G-L+L}{L} \\ d &\geq \frac{G}{L} \end{aligned}$$

That is, when $d \geq \frac{G}{L} + 1$, G is optimal. As such, we want to pay at most $(2 - \frac{G}{L})G = 2G - L$ dollars in total. This is precisely the amount we pay.

Thus, we spend at most $(2 - \frac{G}{L})$ times the minimum amount we would have had to spend if we knew which day the bet would end, as desired. ✓

2. Prove that the above bound is optimal; i.e., any deterministic strategy that you choose cannot result in a ratio less than $(2 - \frac{L}{G})$ in the above bound.

Solution: Let $\frac{L}{G} = \frac{1}{2}$. Consider any deterministic strategy:

- Say the deterministic strategy gives up and pays G immediately. Then imagine the fraud stops the bet on day 1 in which case the optimal strategy was to just pay L on the first day. So the amount we pay is at least $L + G = 2L$, which is $(2 - \frac{L}{G})$ times the optimal amount.
- Say the deterministic strategy pays L until day $d \geq 1$ and then gives up and pays G on day $d+1$. Now suppose the fraud never stops the bet in which case the optimal strategy was to just immediately pay G . So the amount we pay is $d \cdot L + G \geq L + G = 2L$.

So the amount we pay in this case is at least $\frac{2L}{G} = \frac{4L}{2G} = \frac{4L}{L+G} = (2 - \frac{L}{G})$ times the optimal amount. So in either case we cannot achieve a ratio better than $(2 - \frac{L}{G})$. ■

need to
generalize
this argu-
ment to all
values of L
and G .

8

2. Let us now switch to finding an approximate MST. Suppose every edge $e = (u, v) \in E$ of the graph G also has an integer weight $w(e) > 1$ which is known only to players as vertices u and v . The players are also all given a parameter $\epsilon > 0$. They should each send a message of size $\text{poly}((1/\epsilon) \log n)$ to the referee. The referee should then output a spanning tree T such that weight of T is at most $(1 + \epsilon)$ times the weight of the MST of G .

(12.5 points)

Hint: The approach in Lecture 14 was to implement Boruvka's algorithm by finding any edge out of each connected vertex. Can you generalize the approach to find a $(1 + \epsilon)$ -approximate MST instead of a single edge instead? You should then be able to run Boruvka's algorithm to find an approximate MST not yet seen.

Solution: Since we're looking for an $(1 + \epsilon)$ -approximate MST of G , each time we add an edge to the spanning tree, we want the difference between the weight of this edge and the weight of the edge that we're supposed to pick for a MST at this time to be within a factor of $(1 + \epsilon)$. We can also think of it as we're trying to find a minimum spanning tree where edges have weights scaled by a factor of $(1 + \epsilon)$. The weight of each edge in the MST of G will be at most a factor of $(1 + \epsilon)$ times the weight of the MST of G . In this sense, we can group edges of weights within the range of $[1 + \epsilon^j, (1 + \epsilon)^{j+1})$ as Type j for $j \in [0, \log_{1+\epsilon}(\text{poly}(n))]$ since the weight of each edge in the MST of G will be at most a factor of $(1 + \epsilon)$ times the weight of the MST of G .

Algorithm:

- Let $k = \log_{1+\epsilon}(\text{poly}(n))$ and $t = \log n$. The players and referee sample $t - k$ matrices A_i where $i \in [0, t]$ and A_i is such that $A_i A_i^\top = D$ described in Lecture 14.
- For each $i \in [0, t]$, the referee samples P_i for $v \in V$ such that $A_i(v)$ is a vector similar to the one introduced in claim 1 in box 14 except that $x_i(v)$ is designed for edges incident on v with $w(v) \in [1 + \epsilon^j, (1 + \epsilon)^{j+1})$.
- The referee then builds a spanning tree for each sub-graph corresponding to each weight range $w(v) \in [1 + \epsilon^j, (1 + \epsilon)^{j+1})$. The referee first builds the connected component using the edge in the first weight range ($j = 1$). Then the referee computes the sketch of this new component when considering the edges in the second weight range ($j = 2$) and move on until all vertices are covered.

why does this algorithm successfully find an approximate MST? what is the number of bits?

Q5 0 / 15

Refer to the assignment
PDF.

This question wasn't answered