

Lecture 10

October 16, 2023

Instructor: Sepehr Assadi

Disclaimer: *These notes have not been subjected to the usual scrutiny reserved for formal publications. They may be distributed outside this class only with the permission of the Instructor.*

Topics of this Lecture

- | | | |
|---|-----------------------------------|---|
| 1 | The Set Cover Problem | 1 |
| 2 | Randomized Rounding for Set Cover | 2 |
| 3 | Dual Fitting for Set Cover | 3 |

1 The Set Cover Problem

We continue our study of **approximation algorithms** via **linear programming and bounding integrality gaps**. In this lecture, we will see two different examples of approximation algorithms, both for the (weighted) set cover problem.

Recall that in the **set cover problem**, we have a collection of m sets S_1, \dots, S_m from a universe U of n elements. We further assume that each set S has a weight $w(S) > 0$. The goal is to find the minimum weight collection of sets that *cover* the universe U , i.e.,

$$\min_{T \subseteq [m]} \sum_{i \in T} w(S_i) \quad \text{such that} \quad \bigcup_{i \in T} S_i = U.$$

This is a very famous **NP-hard problem** with tons of applications. It also generalizes various problems, including the **minimum vertex cover**, we saw in the last lecture.

The following is an LP relaxation of the set cover problem.

$$\begin{aligned} & \min_{x \in \mathbb{R}^m} \quad \sum_S w(S) \cdot x_S \\ & \text{subject to} \quad \sum_{S \ni e} x_S \geq 1 \quad \forall e \in U, \\ & \quad \quad \quad x_S \geq 0 \quad \forall S. \end{aligned}$$

To see why this is a relaxation, consider the case when $x_S \in \{0, 1\}$ for all S with the interpretation that $x_S = 1$ iff we pick S in the solution; then, the objective is minimizing the weight of the solution and each constraint ensures that we pick at least one set that covers this particular element in the solution.

Recall that in the previous lecture, we considered a **deterministic rounding scheme** for the minimum vertex cover problem. We can do a similar type of argument for the set cover LP also. However, the bounds we get are going to be quite weak in general and depend on the maximum number of sets that cover any fixed element (basically the maximum “length” of the summation in the constraints of the above LP). You are encouraged to work out the details of this rounding scheme on your own.

2 Randomized Rounding for Set Cover

We now consider another simple rounding scheme for the set cover LP, this time by randomly rounding each set to 1 depending on the value the LP assigned to it.

In particular, suppose we solve the set cover LP and obtain a *fractional* solution $x \in \mathbb{R}^m$. Then, it is natural to think that the sets S with “large” values of x_S are “more important” than the ones with lower values of x_S . So, if we want to pick a set to include in our solution, we may want to prioritize picking sets with larger values of x_S over the smaller ones. However, instead of doing this deterministically, we are going to do it *randomly*, namely, sample each set S with a probability p_S proportional to x_S , i.e.,

$$p_S \propto x_S.$$

The question is then what to pick for the *scaling factor*? Let us pick a parameter $\beta \geq 1$ for this, i.e., sample each set S with probability $p_S := \beta \cdot x_S$ (where x_S is the solution obtained from the LP). I.e.

Algorithm 1. A randomized rounding algorithm for set cover.

1. Solve the LP relaxation of the set cover problem to obtain the optimal fractional solution $x \in \mathbb{R}^m$.
2. For every set S , add S to the final solution, denoted by ALG , with probability

$$p_S := \min(\beta \cdot x_S, 1)$$

where β is a parameter to be determined later^a.

^aIf you really cannot wait for later, β is going to be $2 \ln n$...

The analysis of this algorithm is in two steps. We first show that we can always upper bound the expected approximation ratio of the algorithm for any choice of β . In the following, we use $w(ALG)$ to denote the weight of the sets in ALG , i.e., the weight of the solution returned by the algorithm.

Claim 1. $\mathbb{E}[w(ALG)] = \beta \cdot \text{opt}_{LP}$, where opt_{LP} is the optimum objective value of the LP.

Proof. We have,

$$\begin{aligned} \mathbb{E}[w(ALG)] &= \mathbb{E}\left[\sum_S \mathbb{I}(S \in ALG) \cdot w(S)\right] && \text{(because we pay a weight of } w(S) \text{ for every set } S \in ALG) \\ &= \sum_S w(S) \mathbb{E}[\mathbb{I}(S \in ALG)] && \text{(by linearity of expectation)} \\ &= \sum_S w(S) \Pr(S \in ALG) && \text{(as the expected value of an indicator random variable is its probability of being one)} \\ &= \sum_S w(S) \cdot p_S && \text{(by the definition of } p_S) \\ &= \sum_S w(S) \cdot \beta \cdot x_S && \text{(by the choice of } \beta) \\ &= \beta \cdot \text{opt}_{LP} && \text{(by the definition of the objective value of the LP)} \end{aligned}$$

concluding the proof. \square

Claim 1 ensures that we can always bound the expected weight of our solution in terms of the optimal LP for any choice of β that we decide to pick. We now establish a lower bound on the probability that ALG output by the algorithm is an actual set cover, i.e., is a feasible solution.

Claim 2. $\Pr(\text{ALG is a \textit{not} a feasible set cover}) \leq n \cdot e^{-\beta}.$

Proof. By union bound,

$$\begin{aligned} \Pr(\text{ALG is a \textit{not} a feasible set cover}) &= \Pr(\text{there exists an element not covered by ALG}) \\ &\leq \sum_{e \in U} \Pr(\text{ALG does not cover } e). \end{aligned}$$

Fix any $e \in U$ and let us bound the probability term for e in the RHS above. Note that we can assume without loss of generality that for every S that covers e , $p_S < 1$ as if $p_S = 1$, we will deterministically cover e and so this probability is zero. We have,

$$\begin{aligned} \Pr(\text{ALG does not cover } e) &= \prod_{S \ni e} (1 - p_S) \\ &\quad \text{(none of the sets } S \ni e \text{ should be picked, and the choices are independent)} \\ &\leq \prod_{S \ni e} \exp(-p_S) \quad \text{(as } 1 - x \leq e^{-x} \text{ for all } x \in (0, 1)) \\ &= \exp\left(-\sum_{S \ni e} p_S\right) \quad \text{(as } e^a \cdot e^b = e^{a+b} \text{ for all } a, b) \\ &= \exp\left(-\sum_{S \ni e} \beta \cdot x_S\right) \quad \text{(as we have } p_S = \beta \cdot x_S \text{ by our assumption that } p_S < 1) \\ &\leq \exp(-\beta \cdot 1) \quad \text{(as } x \text{ is a feasible solution to the LP and thus } \sum_{S \ni e} x_S \geq 1) \\ &= e^{-\beta}. \end{aligned}$$

The bound in the claim now follows as there are n elements in total. \square

Claim 2 now implies that if we pick $\beta = 2 \ln n$, then, with probability $1 - 1/n$, we obtain a feasible set cover as well (and that if $\beta = o(\ln n)$, we are really not going to get a feasible cover – prove this for yourself).

Finally, we can also apply Markov bound to Claim 1 and say that for $\beta = 2 \ln n$,

$$\Pr(w(\text{ALG}) > 8 \ln n \cdot \text{opt}_{LP}) \leq 1/4.$$

So, by union bound, with probability $1 - (1/4 + 1/n) \geq 2/3$, we obtain a feasible set cover which is an $O(\log n)$ approximation. This concludes our randomized rounding algorithm for set cover.

3 Dual Fitting for Set Cover

We now consider another, less direct, approach for using linear programming to obtain an approximation algorithm for set cover. Instead of first solving the LP and then rounding it, we are going to design a natural greedy algorithm for set cover that on the surface has nothing to do with LPs. But, to analyze this algorithm, we will rely on the notion of *LP duality* that we introduced in the previous lecture.

To continue, we first state, without proof, the following LP as the dual of the set cover LP (you are encouraged to prove this is indeed the dual):

$$\begin{aligned} &\max_{y \in \mathbb{R}^n} \quad \sum_{e \in U} y_e \\ \text{subject to} \quad &\sum_{e \in S} y_e \leq w(S) \quad \forall S, \\ &y_e \geq 0 \quad \forall e. \end{aligned}$$

Let us now consider the following combinatorial algorithm.

Algorithm 2. A greedy algorithm for set cover

1. Let $L = U$ originally be the elements *left* to cover and $ALG = \emptyset$ be the final set cover.
2. While $L \neq \emptyset$
 - (a) Let S be any set that minimizes
$$\frac{w(S)}{|L \cap S|},$$

i.e., the ratio of weight of S to the number of new elements it covers.
 - (b) Add S to ALG and update $L \leftarrow L \setminus S$.
3. Return ALG as a set cover.

The fact that ALG is indeed a feasible set cover is immediate to see. Thus, we only need to bound the approximation value of this algorithm. To do, we are going to use the dual LP. We first define an *infeasible* dual “solution” y' that we can easily relate the cost of our algorithm to, and then show that with a proper scaling down, we can turn y' into a feasible dual solution y . The final bound then is obtained by applying weak duality theorem.

We define $y' \in \mathbb{R}^n$ as follows. In the algorithm, whenever an element e is covered *for the first time* by a set S , i.e., the iteration that we remove e from L , we define

$$y'_e := \frac{w(S)}{|L \cap S|},$$

where L is the set of elements before we pick S in ALG .

Claim 3. $w(ALG) = \sum_{e \in U} y'_e$.

Proof. We have,

$$\begin{aligned}
\sum_{e \in U} y'_e &= \sum_{S \in ALG} \sum_{\substack{e \text{ covered} \\ \text{first by } S}} y'_e && \text{(every element is covered exactly once for the ‘first time’ in } ALG) \\
&= \sum_{S \in ALG} \sum_{\substack{e \text{ covered} \\ \text{first by } S}} \frac{w(S)}{|L \cap S|} \\
&\text{(by definition of } y'_e \text{ where } L \text{ here refers to the set of remaining elements before picking } S \text{ in } ALG) \\
&= \sum_{S \in ALG} |S \cap L| \cdot \frac{w(S)}{|L \cap S|} \\
&\hspace{15em} \text{(because the number of elements covered by } S \text{ for the first time is } |L \cap S|) \\
&= \sum_{S \in ALG} w(S) = w(ALG),
\end{aligned}$$

concluding the proof. □

So, we can relate the cost of our algorithm to the infeasible dual solutions y' . We now show that a scaled down version of y' actually leads to a feasible dual solution.

Claim 4. Define $y = \frac{y'}{\ln(n)}$. Then, y is a feasible dual solution.

Proof. To prove y is a feasible solution, we need to show that for every set S ,

$$\sum_{e \in S} y_e \leq w(S).$$

Let $S = \{e_1, e_2, \dots, e_k\}$ where we additionally assume that e_1 is covered before (or at the same time as) e_2 , e_2 before (or at the same time as) e_3 , and so on and so forth. When e_1 was covered, the “cost” of set S that we minimize in the algorithm greedily was $w(S)/k$ as all its elements are still in L . Since we are picking a set with minimum “cost” and by the definition of y' , we have

$$y'_{e_1} \leq \frac{w(S)}{k}.$$

Then, when we pick e_2 , we have,

$$y'_{e_2} \leq \frac{w(S)}{k-1},$$

because again S can cover at least $k-1$ elements and thus the “cost” of the set we pick is at most the RHS above. Continuing like this, for every $i \in [k]$, we have,

$$y'_{e_i} \leq \frac{w(S)}{k-i+1}.$$

Thus,

$$\sum_{i=1}^k y'_{e_i} \leq \sum_{i=1}^k \frac{w(S)}{k-i+1} = w(S) \cdot \sum_{i=1}^k \frac{1}{i} \leq w(S) \cdot \ln k \leq w(S) \cdot \ln n,$$

where the first inequality is by the sum of Harmonic series and the second is because $k \leq n$. By the definition of $y = y' / \ln(n)$, we obtain that $\sum_e y_e \leq w(S)$ as desired, concluding the proof. \square

To conclude, we have,

$$\begin{aligned} w(ALG) &= \sum_e y'_e && \text{(by Claim 3)} \\ &= \ln(n) \cdot \sum_e y_e && \text{(by the definition of } y) \\ &\leq \ln(n) \cdot \text{opt}_{dual} \\ &\text{(because } y \text{ is a feasible solution to dual LP by Claim 4 which is a maximization LP)} \\ &\leq \ln(n) \cdot \text{opt}_{LP}. && \text{(by weak duality theorem)} \end{aligned}$$

This concludes the proof as opt_{LP} is at most that of the (integral) set cover.

Thus, we obtained a deterministic $\ln(n)$ approximation algorithm to weighted set cover this way.

Remark. In general, a dual fitting approach to a minimization LP works by starting with a basic (combinatorial) algorithm for the problem and using LPs only in its analysis. Using the dual LP of the problem, we show that the primal (integral) solution we picked is “paid for” by some *infeasible* dual. By fully paid for we mean that the objective function value of the primal solution found is at most the objective function value of the dual computed. The main step in the analysis consists of scaling the dual by a suitable factor and showing that the shrunk/expanded dual is now indeed feasible. This dual solution then provides a bound on the optimal solution, and the factor in this bound is the approximation guarantee of the algorithm.